Assignment Date	23 september 2022
Student Name	SUDARSHAN.B
Student Roll Number	311419205043
Maximum Marks	2 Marks

Data Visualization and Pre-processing

Question-1:

- Load the dataset
  Solution:

```
import pandas as pd import
seaborn as sns import
matplotlib.pyplot as plt
importnumpy as np
sns.set_style('darkgrid')
sns.set(font_scale=1.3)
```

In [2]:
```
df=pd.read_excel("/content/Churn_Modelling.xlsx")
```

```
In [1]:   import pandas as pd
          import seaborn as sns
          import matplotlib.pyplot as plt
          import numpy as np
          sns.set_style('darkgrid')
          sns.set(font_scale=1.3)

In [2]:   df=pd.read_excel("/content/Churn_Modelling.xlsx")
```

Question-2:

- Perform
  Below
  Visualizations.
- Univariate Analysis

- Bi - Variate Analysis ●
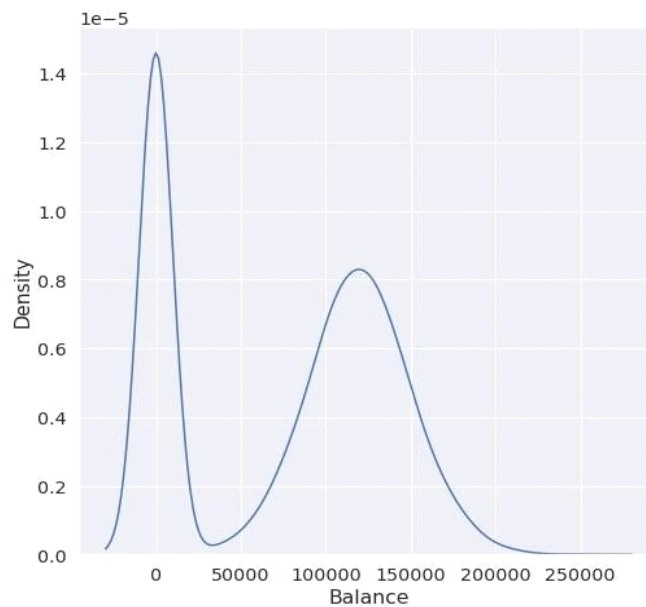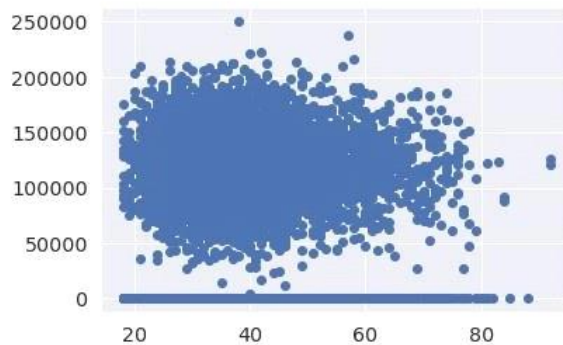 Multi - Variate
  Analysis

Solution:

.

```
#Perform Univariate
Analysis
plt.figure(figsize=(8,8))
sns.kdeplot(x=df['Balance']
)
```

```
In [7]:   #Perform Univariate Analysis
          plt.figure(figsize=(8,8))
          sns.kdeplot(x=df['Balance'])

Out[7]:   <matplotlib.axes._subplots.AxesSubplot at 0x7fc3f3579c50>
```



```
#Perform Bivariate Analysis
plt.scatter(df.Age,df.Balance)
```



```
#Perform Bivariate Analysis df.corr()
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
CreditScore	1.000000	-0.003965	0.000842	0.006268	0.012238	-0.005458	0.025651	-0.001384	-0.027094
Age	-0.003965	1.000000	-0.009997	0.028308	-0.030680	-0.011721	0.085472	-0.007201	0.285323
Tenure	0.000842	-0.009997	1.000000	-0.012254	0.013444	0.022583	-0.028362	0.007784	-0.014001
Balance	0.006268	0.028308	-0.012254	1.000000	-0.304180	-0.014858	-0.010084	0.012797	0.118533
NumOfProducts	0.012238	-0.030680	0.013444	-0.304180	1.000000	0.003183	0.009612	0.014204	-0.047820
HasCrCard	-0.005458	-0.011721	0.022583	-0.014858	0.003183	1.000000	-0.011866	-0.009933	-0.007138
IsActiveMember	0.025651	0.085472	-0.028362	-0.010084	0.009612	-0.011866	1.000000	-0.011421	-0.156128
EstimatedSalary	-0.001384	-0.007201	0.007784	0.012797	0.014204	-0.009933	-0.011421	1.000000	0.012097
Exited	-0.027094	0.285323	-0.014001	0.118533	-0.047820	-0.007138	-0.156128	0.012097	1.000000

```
#Perform Bivariate Analysis
import statsmodels.api as sm

#define response variable y
= df['CreditScore']

#define explanatory variable x
= df[['EstimatedSalary']]

#add constant to predictor variables x
= sm.add_constant(x)

#fit linear regression model model
= sm.OLS(y, x).fit()

#view model summary print(model.summary())
```

```
                          OLS Regression Results
==============================================================================
Dep. Variable:            CreditScore   R-squared:                       0.000
Model:                            OLS   Adj. R-squared:                 -0.000
Method:                 Least Squares   F-statistic:                   0.01916
Date:                Thu, 29 Sep 2022   Prob (F-statistic):              0.890
Time:                        14:58:55   Log-Likelihood:                -59900.
No. Observations:               10000   AIC:                         1.198e+05
Df Residuals:                    9998   BIC:                         1.198e+05
Df Model:                           1
Covariance Type:            nonrobust
==============================================================================
                   coef    std err          t      P>|t|      [0.025      0.975]
------------------------------------------------------------------------------
const            650.7617      1.940    335.407      0.000     646.958     654.565
EstimatedSalary -2.326e-06   1.68e-05     -0.138      0.890   -3.53e-05    3.06e-05
==============================================================================
Omnibus:                      132.939   Durbin-Watson:                   2.014
Prob(Omnibus):                  0.000   Jarque-Bera (JB):               84.242
Skew:                          -0.072   Prob(JB):                     5.10e-19
Kurtosis:                       2.574   Cond. No.                     2.32e+05
==============================================================================

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 2.32e+05. This might indicate that there are
strong multicollinearity or other numerical problems.
```
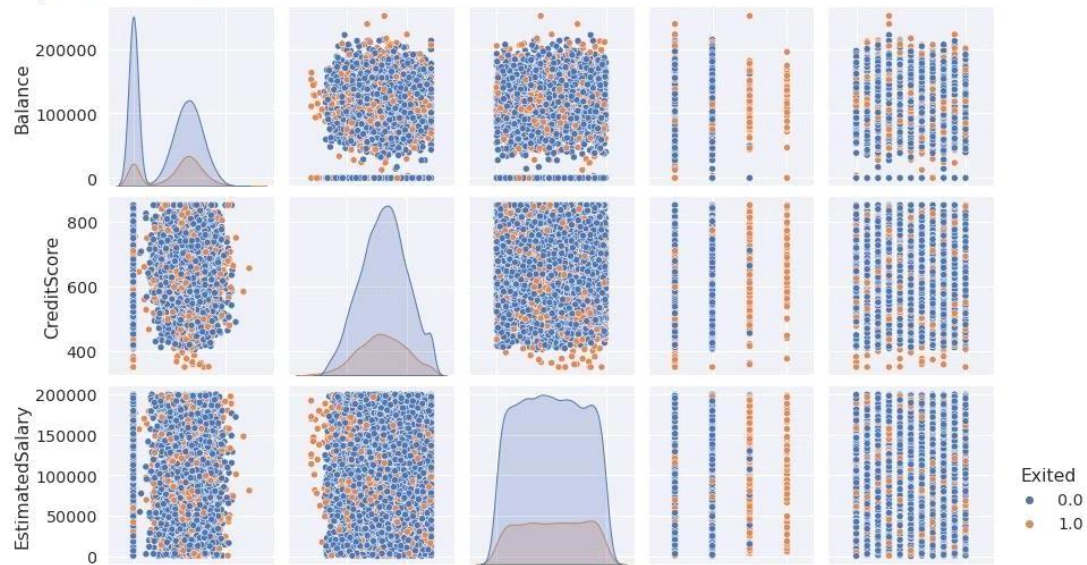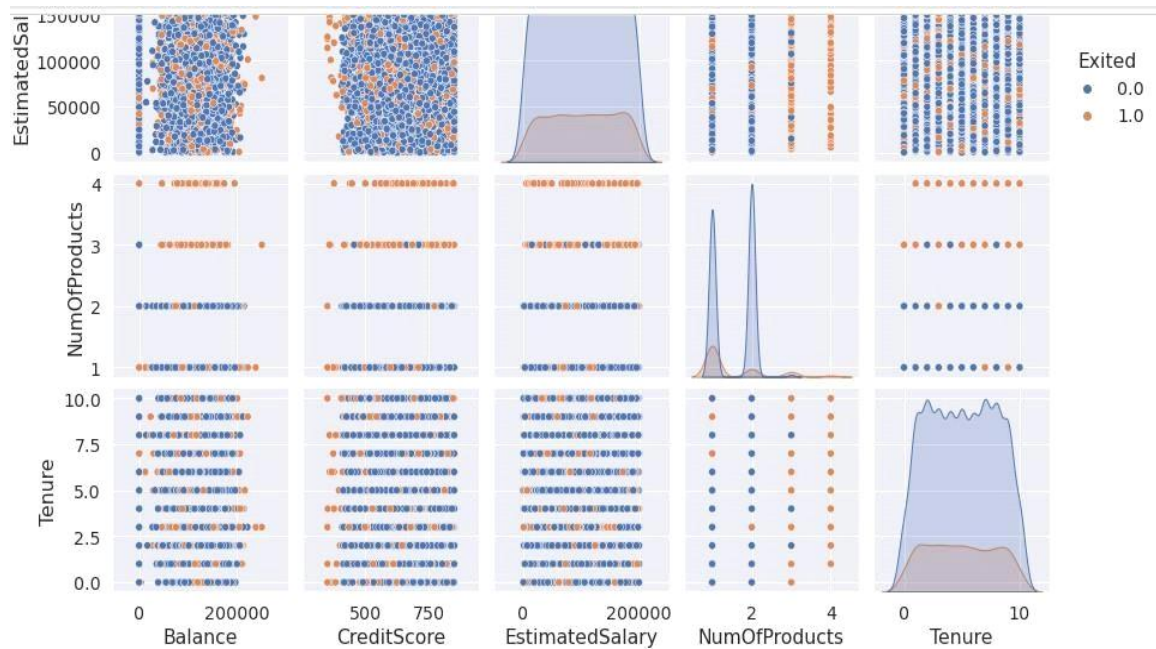
```
#Perform Multivariate Analysis plt.figure(figsize=(4,4))
sns.pairplot(data=df[["Balance","CreditScore","EstimatedSalary","NumOfProdu
cts","Tenure","Exited"]],hue="Exited")
```


`<Figure size 288x288 with 0 Axes>`

Question-3:

- Perform descriptive statistics on the dataset.

Solution:

```
#Perform Descriptive Statistics
df=pd.DataFrame(df) print(df.sum())
```

```
CreditScore                                        6505288.0
Geography          FranceSpainFranceFranceSpainSpainFranceGermany...
Gender             FemaleFemaleFemaleFemaleFemaleMaleMaleFemaleMa...
Age                                                 389218.0
Tenure                                               50128.0
Balance                                          764858892.88
NumOfProducts                                        15302.0
HasCrCard                                             7055.0
IsActiveMember                                        5151.0
EstimatedSalary                                  1000902398.81
Exited                                                2037.0
dtype: object
```

```
#Perform Descriptive Statistics print("----
Sum Value-----") print(df.sum(1)) print("---
```

```
                                        ") print("
--Product Value-----") print(df.prod())
print("
                                                        "
)
```

```
----Sum Value-----
0        102015.88
1        197002.44
2        274149.37
3         94567.63
4        205492.92
            ...
9995      97088.64
9996     159633.38
9997      42840.58
9998     168784.83
9999     169159.57
Length: 10000, dtype: float64
---------------------------------
-----Product Value-----
CreditScore        inf
Age                inf
Tenure             0.0
Balance            0.0
NumOfProducts      inf
HasCrCard          0.0
IsActiveMember     0.0
EstimatedSalary    inf
Exited             0.0
dtype: float64
---------------------------------
```

```
#Perform Descriptive Statistics print("---------
-Mean Value-----------")
print(df.mean())print("
") print("---------Median Value----- ")
print(df.median()) print("---------------------
        ") print("            Mode Value
---") print(df.mode())
print("                                    ")
```

```
----------Mean Value-----------
CreditScore          650.528800
Age                   38.921800
Tenure                 5.012800
Balance            76485.889288
NumOfProducts          1.530200
HasCrCard              0.705500
IsActiveMember         0.515100
EstimatedSalary   100090.239881
Exited                 0.203700
dtype: float64
-----------------------------------
----------Median Value---------
CreditScore          652.000
Age                   37.000
Tenure                 5.000
Balance            97198.540
NumOfProducts          1.000
HasCrCard              1.000
IsActiveMember         1.000
EstimatedSalary   100193.915
Exited                 0.000
dtype: float64
-----------------------------------
----------Mode Value------------
   CreditScore Geography Gender   Age   Tenure   Balance   NumOfProducts
0       850.0    France    Male  37.0      2.0       0.0             1.0

   HasCrCard   IsActiveMember   EstimatedSalary   Exited
0        1.0              1.0          24924.92      0.0
-----------------------------------
```

Question-4:

• Handle the Missing values

Solution:

#Handling with missing Values df.isnull().values;
#Checking values are null

#Handling with missing Values
df.notnull()*#Checking values are not null*

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	True	True	True	True	True	True	True	True	True	True	True
1	True	True	True	True	True	True	True	True	True	True	True
2	True	True	True	True	True	True	True	True	True	True	True
3	True	True	True	True	True	True	True	True	True	True	True
4	True	True	True	True	True	True	True	True	True	True	True
...	...	...	...	...	...	...	...	...	...	...	...
9995	True	True	True	True	True	True	True	True	True	True	True
9996	True	True	True	True	True	True	True	True	True	True	True
9997	True	True	True	True	True	True	True	True	True	True	True
9998	True	True	True	True	True	True	True	True	True	True	True
9999	True	True	True	True	True	True	True	True	True	True	True

10000 rows × 11 columns

Question-5:

• Find the outliers and replace the outliers

Solution:

```
#Find outliers & replace the outliers
sns.boxplot(df['Balance'])
```

```
#Find outliers & replace the outliers
print(np.where(df['Balance']>100000))

(array([   2,    4,    5, ..., 9987, 9993, 9999]),)
```

```
#Find outliers & replace the outliersfrom scipy import stats
import numpyas np

z = np.abs(stats.zscore(df["EstimatedSalary"]))print(z)
```

In [19]:

```
0         0.021886
1         0.216534
2         0.240687
3         0.108918
4         0.365276
           ...
9995      0.066419
9996      0.027988
9997      1.008643
9998      0.125231
9999      1.076370
Name: EstimatedSalary, Length: 10000, dtype: float64
```

Question-6:

• Check for Categorical columns and perform encoding

Solution:

```
#Check for categorical columns & performs encoding from
sklearn.preprocessing import LabelEncoder

df['Gender'].unique()
df['Gender'].value_counts()
encoding=LabelEncoder()
df["Gender"]=encoding.fit_transform(df.iloc[:,1].valu
es) df
```

```
]:  #Check for categorical columns & performs encoding
    from sklearn.preprocessing import LabelEncoder
    df['Gender'].unique()

]:  array(['Female', 'Male'], dtype=object)
```

```
]:  #Check for categorical columns & performs encoding
    df['Gender'].value_counts()

]:  Male      5457
    Female    4543
    Name: Gender, dtype: int64
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	619.0	France	0	42.0	2.0	0.00	1.0	1.0	1.0	101348.88	1.0
1	608.0	Spain	2	41.0	1.0	83807.86	1.0	0.0	1.0	112542.58	0.0
2	502.0	France	0	42.0	8.0	159660.80	3.0	1.0	0.0	113931.57	1.0
3	699.0	France	0	39.0	1.0	0.00	2.0	0.0	0.0	93826.63	0.0
4	850.0	Spain	2	43.0	2.0	125510.82	1.0	1.0	1.0	79084.10	0.0
...	...	...	...	...	...	...	...	...	...	...	...
9995	771.0	France	0	39.0	5.0	0.00	2.0	1.0	0.0	96270.64	0.0
9996	516.0	France	0	35.0	10.0	57369.61	1.0	1.0	1.0	101699.77	0.0
9997	709.0	France	0	36.0	7.0	0.00	1.0	0.0	1.0	42085.58	1.0
9998	772.0	Germany	1	42.0	3.0	75075.31	2.0	1.0	0.0	92888.52	1.0
9999	792.0	France	0	28.0	4.0	130142.79	1.0	1.0	0.0	38190.78	0.0

10000 rows × 11 columns

Question-7:

• Split the data into dependent and independent variables.

Solution:

```
#Split the data into Dependent & Independent Variables
print("----------Dependent Variables--------- ")
X=df.iloc[:,1:4] print(X) print("--------------------
-
                  ") print("          Independent
Variables---------") Y=df.iloc[:,4] print(Y)
print("

                                              "
)
```

Question-8:

• Scale the independent variables

Solution:

```
        #Split the data into Dependent & Independent Variables
print("----------Dependent Variables--------- ")
X=df.iloc[:,1:4] print(X) print("---------------
                  ") print("
Independent Variables---------")
Y=df.iloc[:,4]print(Y)
print("                                        ")
```

Question-9:

• Split the data into training and testing

Solution:

```
#Split the data into training & testing from
sklearn.model_selection import train_test_split
```

```
#Split the data into training & testing
x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=4,random_state=4) x_train x_test y_train y_test
```

In [34]:

Out[31]:

	const	EstimatedSalary
1603	1.0	23305.85
8713	1.0	41248.80
4561	1.0	143317.42
6600	1.0	174123.16

```
In [32]:  #Split the data into training & testing
          y_train
```

```
Out[32]:  2558    727.0
          7642    811.0
          8912    623.0
          3319    430.0
          6852    600.0
                  ...
          456     733.0
          6017    487.0
          709     686.0
          8366    637.0
          1146    614.0
          Name: CreditScore, Length: 9996, dtype: float64
```

	const	EstimatedSalary
2558	1.0	137903.54
7642	1.0	121765.00
8912	1.0	109470.34
3319	1.0	2923.61
6852	1.0	7312.25
...	...	...
456	1.0	7666.73
6017	1.0	9085.00
709	1.0	147794.63
8366	1.0	102515.42
1146	1.0	54776.64

9996 rows × 2 columns